

Looking for Real Exam Questions for IT Certification Exams!

We guarantee you can pass any IT certification exam at your first attempt with just 10-12 hours study of our guides.

Our study guides contain actual exam questions; accurate answers with detailed explanation verified by experts and all graphics and drag-n-drop exhibits shown just as on the real test.

To test the quality of our guides, you can download the one-fourth portion of any guide from <http://www.certificationking.com> absolutely free. You can also download the guides for retired exams that you might have taken in the past.

For pricing and placing order, please visit <http://certificationking.com/order.html>

We accept all major credit cards through www.paypal.com

For other payment options and any further query, feel free to mail us at info@certificationking.com

QUESTION NO: 1

Which of the following describe the functions of a scheduling algorithm? (Choose 4)

- A. Reduce the total amount of computation necessary to complete a job.
- B. Allow short Jobs to complete even when large, long jobs (consuming a lot of resources) are running.
- C. Support the implementation of service-level agreements for multiple cluster users.
- D. Allow multiple users to share clusters in a predictable, policy-guided manner.
- E. Run jobs at periodic times of the day.
- F. Reduce job latencies in an environment with multiple jobs of different sizes.

Answer: A,D,E,F

Explanation:

QUESTION NO: 2

You are running a Hadoop cluster with all monitoring facilities properly configured. Which scenario will go undetected?

- A. Map or reduce tasks that are stuck in an infinite loop.
- B. HDFS is almost full.
- C. The NameNode goes down.
- D. A DataNode is disconnected from the cluster.
- E. MapReduce jobs that are causing excessive memory swaps.

Answer: A

Explanation:

QUESTION NO: 3

Which of the following scenarios makes HDFS unavailable?

- A. JobTracker failure
- B. TaskTracker failure
- C. DataNode failure
- D. NameNode failure
- E. Secondary NameNode failure

Answer: C

Reference:<http://stackoverflow.com/questions/12362727/when-will-hdfs-be-unavailable>

QUESTION NO: 4

What's the relationship between JobTrackers and TaskTrackers?

- A.** The JobTracker runs on a single master node and accepts MapReduce jobs from clients. A TaskTracker runs on every slave node and is responsible for managing actual map and reduce tasks.
- B.** Every node in the cluster runs both a JobTracker and a TaskTracker. The JobTrackers manage jobs, and the TaskTrackers are responsible for managing actual map and reduce tasks.
- C.** The TaskTrackers run on a single master node and accept MapReduce jobs from clients. A JobTracker runs on every slave node and is responsible for managing map and reduce tasks.
- D.** The JobTracker runs on a single master node, but for each separate instance of itself for every client MapReduce job. A TaskTracker runs on every slave node and is responsible for managing actual map and reduce tasks.

Answer: B

Reference:http://hadoop.apache.org/mapreduce/docs/r0.22.0/mapred_tutorial.html (Overview, 4th paragraph)

QUESTION NO: 5

Assuming a large properly configured multi-rack Hadoop cluster, which scenario should not result in loss of HDFS data assuming the default replication factor settings?

- A.** Ten percent of DataNodes simultaneously fail.
- B.** All DataNodes simultaneously fail.
- C.** An entire rack fails.
- D.** Multiple racks simultaneously fail.
- E.** Seventy percent of DataNodes simultaneously fail.

Answer: A

Reference:<http://stackoverflow.com/questions/12399197/in-a-large-properly-configured-multi-rack-hadoop-cluster-which-scenarios-will-b>

QUESTION NO: 6

Which daemon spawns child JVMs to perform MapReduce processing?

- A. JobTracker
- B. NameNode
- C. DataNode
- D. TaskTracker
- E. Secondary NameNode

Answer: D

Reference:[http://www.mindmeister.com/75831919/hadoop-talk-nathan-milford-outbrain\(search Task Tracker\)](http://www.mindmeister.com/75831919/hadoop-talk-nathan-milford-outbrain(search%20Task%20Tracker))

QUESTION NO: 7

A client wants to read a file from HDFS. How does the data get from the DataNodes to the client?

- A. The NameNode reads the blocks from the DataNodes, and caches them. Then, the application reads the blocks from the NameNode.
- B. The application reads the blocks directly from the DataNodes.
- C. The blocks are sent to a single DataNode, then the application reads the blocks from that Data Node.

Answer: B

Reference:<http://kazman.shidler.hawaii.edu/ArchDocOverview.html>

QUESTION NO: 8

What would be a reasonable configuration of disk drives in a Hadoop datanode?

- A. Four 1TB disk drives in a RAID configuration
- B. One 1TB disk drive
- C. Four 1TB disk drives in a JBOD configuration